

CodeForces problems classification

By Ahmed Bouali

ENV SETUP

```
!unzip code_classification_dataset.zip
!pip install pytorch-lightning

inflatng: code_classification_dataset/sample_2330.json
inflatng: code_classification_dataset/sample_2335.json
inflatng: code_classification_dataset/sample_2339.json
inflatng: code_classification_dataset/sample_2344.json
inflatng: code_classification_dataset/sample_2349.json
inflatng: code_classification_dataset/sample_2352.json
inflatng: code_classification_dataset/sample_2361.json
inflatng: code_classification_dataset/sample_2364.json
inflatng: code_classification_dataset/sample_2376.json
inflatng: code_classification_dataset/sample_2377.json
inflatng: code_classification_dataset/sample_2378.json
inflatng: code_classification_dataset/sample_2382.json
inflatng: code_classification_dataset/sample_2383.json
inflatng: code_classification_dataset/sample_2384.json
inflatng: code_classification_dataset/sample_2387.json
inflatng: code_classification_dataset/sample_2391.json
inflatng: code_classification_dataset/sample_2392.json
inflatng: code_classification_dataset/sample_2398.json
inflatng: code_classification_dataset/sample_2414.json
inflatng: code_classification_dataset/sample_2415.json
inflatng: code_classification_dataset/sample_2416.json
inflatng: code_classification_dataset/sample_2426.json
inflatng: code_classification_dataset/sample_2430.json
inflatng: code_classification_dataset/sample_2435.json
inflatng: code_classification_dataset/sample_2438.json
inflatng: code_classification_dataset/sample_2440.json
inflatng: code_classification_dataset/sample_2447.json
inflatng: code_classification_dataset/sample_2449.json
inflatng: code_classification_dataset/sample_2456.json
inflatng: code_classification_dataset/sample_2457.json
inflatng: code_classification_dataset/sample_2465.json
inflatng: code_classification_dataset/sample_2485.json
inflatng: code_classification_dataset/sample_2486.json
inflatng: code_classification_dataset/sample_2488.json
inflatng: code_classification_dataset/sample_2499.json
inflatng: code_classification_dataset/sample_2506.json
inflatng: code_classification_dataset/sample_2510.json
inflatng: code_classification_dataset/sample_2520.json
inflatng: code_classification_dataset/sample_2522.json
inflatng: code_classification_dataset/sample_2524.json
inflatng: code_classification_dataset/sample_2529.json
inflatng: code_classification_dataset/sample_2531.json
inflatng: code_classification_dataset/sample_2538.json
inflatng: code_classification_dataset/sample_2540.json
inflatng: code_classification_dataset/sample_2557.json
inflatng: code_classification_dataset/sample_2565.json
inflatng: code_classification_dataset/sample_2571.json
inflatng: code_classification_dataset/sample_2572.json
inflatng: code_classification_dataset/sample_2573.json
inflatng: code_classification_dataset/sample_2585.json
inflatng: code_classification_dataset/sample_2592.json
inflatng: code_classification_dataset/sample_2596.json
inflatng: code_classification_dataset/sample_2601.json
inflatng: code_classification_dataset/sample_2603.json
inflatng: code_classification_dataset/sample_2607.json
inflatng: code_classification_dataset/sample_2608.json
inflatng: code_classification_dataset/sample_2617.json
inflatng: code_classification_dataset/sample_2622.json
inflatng: code_classification_dataset/sample_2645.json
```

```
import json
import sys
import os

import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import matplotlib.patches as mpatches
import pytorch_lightning as pl
import torch

from time import sleep
from collections import Counter
```



```
'lang': 'Python 3',
'lang_cluster': 'Python',
'difficulty': 1500,
'file_name': 'train_021.jsonl',
'code_uid': 'f1e020a8afc8535f9b870a7df58949d8',
'prob_desc_memory_limit': '256 megabytes',
'prob_desc_sample_inputs': '["7 4\n4727447", "4 2\n4478"]',
'exec_outcome': 'PASSED',
'source_code': "# 447 or 477 start with odd\n\nnn, k = map(int, input().split())\ns = input()\ns = list(s)\n\nfor i in
range(n - 1):\n    if not k:\n        break\n    if s[i] != '4':\n        continue\n    tt = ''.join(s[i:i + 3])\n    if tt in
('447', '477') and i % 2 == 0:\n        if k % 2 == 1 and tt == '447':\n            s[i + 1] = '7'\n            if k % 2 == 1 and
tt == '477':\n                s[i + 1] = '4'\n                break\n            if s[i] == '4' and s[i + 1] == '7':\n                if i % 2 == 0:\n
s[i + 1] = '4'\n            else:\n                s[i] = '7'\n                k -= 1\n\nprint(''.join(s))\n",
'prob_desc_created_at': '1319727600',
'tags': ['strings'],
'hidden_unit_tests': ''}
```

```
df_combined = pd.DataFrame(all_data)
display(df_combined.head())
```

	prob_desc_time_limit	prob_desc_sample_outputs	src_uid	prob_desc_notes	prob_desc_description	prob_desc_created_at
0	2 seconds	["36", "1728"]	9a5bd9f937da55c3d26d5ecde6e50280	Not in the first sample n = 2 · 3 = 6. The divis...	Ayrat has number n, represented as it's prime ...	1319727600
1	2 seconds	["4427477", "4478"]	8ce1ba0a98031c1bc28f53c11905391c	Not in the first sample the number changes in ...	Petya loves lucky numbers. Everybody knows tha...	1319727600
2	2 seconds	["1\n0\n9\n7\n6"]	867b01e7141ef077964a8a0d4c6b762b	Not in the first test case, we can perform one...	You are given a square grid with \$\$\$n\$\$\$ rows ...	1319727600
3	1 second	["0", "1", "2"]	a32db37cb2ebe8945a4c2f32fa2d7fc8	None	Polycarp loves geometric progressions — he col...	1319727600
4	1 second	["2 5", "1 8", "6 9"]	1d4aaf15e5c6fcde50515880aae74720	None	In this problem we consider a special type of ...	1319727600

5 rows × 7 columns

```
df_combined.shape
```

(4982, 7)

```
df_combined.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 4982 entries, 0 to 4981
Data columns (total 21 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   prob_desc_time_limit                  4982 non-null   object
1   prob_desc_sample_outputs              4982 non-null   object
2   src_uid                              4982 non-null   object
3   prob_desc_notes                       3632 non-null   object
4   prob_desc_description                 4982 non-null   object
5   prob_desc_output_spec                 4897 non-null   object
6   prob_desc_input_spec                  4949 non-null   object
7   prob_desc_output_to                   4981 non-null   object
8   prob_desc_input_from                  4981 non-null   object
9   lang                                  4982 non-null   object
10  lang_cluster                          4982 non-null   object
11  difficulty                            4943 non-null   float64
12  file_name                             4982 non-null   object
13  code_uid                              4982 non-null   object
14  prob_desc_memory_limit                4982 non-null   object
15  prob_desc_sample_inputs               4982 non-null   object
16  exec_outcome                          4982 non-null   object
17  source_code                           4982 non-null   object
18  prob_desc_created_at                  4982 non-null   object
19  tags                                  4982 non-null   object
20  hidden_unit_tests                     4982 non-null   object
dtypes: float64(1), object(20)
memory usage: 817.5+ KB
```

```
dup_value_counts = df_combined.drop(columns=['tags']).apply(lambda s: s.value_counts().gt(1).sum())
dup_value_counts.sort_values(ascending=True)
```

	0
src_uid	0
prob_desc_description	0
code_uid	0
source_code	0
exec_outcome	1
hidden_unit_tests	1
lang_cluster	1
prob_desc_output_to	3
prob_desc_input_from	3
lang	5
prob_desc_memory_limit	9
prob_desc_time_limit	18
prob_desc_input_spec	21
difficulty	29
prob_desc_notes	35
prob_desc_output_spec	61
prob_desc_sample_inputs	66
file_name	111
prob_desc_sample_outputs	277
prob_desc_created_at	1109

dtype: int64

df_combined.isna().sum()

	0
prob_desc_time_limit	0
prob_desc_sample_outputs	0
src_uid	0
prob_desc_notes	1350
prob_desc_description	0
prob_desc_output_spec	85
prob_desc_input_spec	33
prob_desc_output_to	1
prob_desc_input_from	1
lang	0
lang_cluster	0
difficulty	39
file_name	0
code_uid	0
prob_desc_memory_limit	0
prob_desc_sample_inputs	0
exec_outcome	0
source_code	0
prob_desc_created_at	0
tags	0
hidden_unit_tests	0

dtype: int64

Most importantly we don't have duplications in **src_uid** and **prob_desc_description** so we can assume that the problems are unique

```
df_combined.difficulty.describe()
```

```

      difficulty
count 4943.000000
mean  1668.336031
std    568.273766
min    -1.000000
25%    1200.000000
50%    1700.000000
75%    2100.000000
max    3500.000000
dtype: float64

```

[illegible]

All DESC columns are informative, should be kept (but we need to deal with the empty vals)

```
df_combined[['lang','lang_cluster']].drop_duplicates()
```

	lang	lang_cluster
0	Python 2	Python
1	Python 3	Python
2	PyPy 3-64	Python
4	PyPy 3	Python
18	PyPy 2	Python

All problems are in python, we can drop lang

- ✓ FEATURE ENGINEERING AND DATASET FILTERING BASED ON GIVEN TAGS

```
data = df_combined.copy()
```

```
#lets keep only relevant cols
data = data[['src_uid', 'prob_desc_notes', 'prob_desc_description', 'prob_desc_output_spec', 'prob_desc_input_spec', 'difficulty', 's
```

```
data.shape, data.dropna().shape
```

 $((4982, 8), (3527, 8))$

```
#We concatenate problem features into one big paragraph
data['problem description'] = data['prob desc description'].fillna("") + " " + data['prob desc notes'].fillna("") + " " + data
```

```
data = data[['src_uid', 'problem_description', 'difficulty', 'source_code', 'tags']]
```

```
# lets drop rows with empty difficulty values (we only lose around 40 exp, less than 1% of the dataset)
data.shape, data.dropna().shape
```

 $((4982, 5), (4943, 5))$

```
data = data.dropna()
```

```
tags = list(set().union(*data.tags.values.tolist()))
```

```
len(tags)
```

```
37
```

```
#How are tags distributed accross dataset ?
tags_occ = Counter()
for row in data.tags.values.tolist():
    tags_occ.update(row)
```

```
len(tags_occ)
```

```
37
```

```
focus_tags = ['math', 'graphs', 'strings', 'number theory', 'trees', 'geometry', 'games', 'probabilities']
```

```
tags_occ_sorted = sorted(tags_occ.items(), key=lambda x: x[1], reverse=True)
```


```
focus_color = 'tomato'
other_color = 'skyblue'
tags_sorted = [tag[0] for tag in tags_occ_sorted]
colors = [focus_color if tag in focus_tags else other_color for tag in tags_sorted]
counts = [tags_occ[tag] for tag in tags_sorted]

plt.figure(figsize=(8, max(4, len(tags) * 0.35)))
y = range(len(tags_sorted))
bars = plt.barh(y, counts, color=colors)
plt.yticks(y, tags_sorted)
plt.xlabel('Count')
plt.title('Tag occurrences highlighted')

for bar, count in zip(bars, counts):
    width = bar.get_width()
    plt.text(width + max(1, 0.01 * max(counts)),
             bar.get_y() + bar.get_height() / 2,
             str(count),
             va='center', ha='left', fontsize=9)

plt.legend([mpatches.Patch(color=focus_color), mpatches.Patch(color=other_color)],
           ['Focus tags', 'Other tags'])

plt.tight_layout()
plt.show()
```



We can see that the 8 chosen classes are not the top 8 classes in the dataset, There's still a huge disparity between the largest class (math : 1409) and the smallest (probabilities : 92)

But let's ask ourselves this question :

can an example be of class probabilities without being of class maths ?

```
data[data.tags.apply(lambda x: 'probabilities' in x and 'math' not in x)][['problem_description', 'tags']].head(5)
```

	problem_description	tags	
282	The Bad Luck Island is inhabited by three kind...	[dp, probabilities]	
530	Dexterina and Womandark have been arch-rivals ...	[games, probabilities, matrices]	
831	Masha lives in a country with \$\$\$n\$\$\$ cities n...	[dp, probabilities, graphs, brute force]	
854	You are given a string \$\$\$s\$\$\$ consisting of \$...	[brute force, dp, greedy, probabilities]	
861	One night, having had a hard day at work, Pety...	[probabilities, sortings, binary search, dfs a...	

```
data[data.tags.apply(lambda x: 'probabilities' in x and 'math' not in x)][['problem_description', 'tags']].shape
```

```
(46, 2)
```

```
tags_occ['probabilities']
```

```
91
```

Now we know that we have a multilabel classification problem with distinct labels. An exercise that is labeled as "probabilities" can be labeled also as "math" and can be labeled only as "probabilities". We have also an imbalanced dataset.

Lets filter the dataset based on the 8 tags

```
data.tags = data.tags.apply(lambda x: list(set(x).intersection(set(focus_tags))))
```

```
data.isna().sum()
```

```

src_uid      0
problem_description  0
difficulty    0
source_code   0
tags          0
dtype: int64
```

```
#Since tags col is a list col, we will check for empty values seperately
data[data.tags.apply(lambda x: len(x) == 0)]
```

	src_uid	problem_description	difficulty	source_code	tags	
2	867b01e7141ef077964a8a0d4c6b762b	You are given a square grid with \$\$\$n\$\$\$ rows ...	1200.0	import sys\n\ninput = lambda: sys.stdin.readli...	[]	
3	a32db37cb2ebe8945a4c2f32fa2d7fc8	Polycarp loves geometric progressions — he col...	2200.0	def main():\n n = int(input())\n l = tup...	[]	
4	1d4aaf15e5c6fcde50515880aee74720	In this problem we consider a special type of ...	800.0	garbage = int(input())\nz = list(map(int, inpu...	[]	
5	91be5db48b44a4adff4c809ffb8e3e	Caisa is going to have a party and he needs to...	1200.0	n,s=map(int,raw_input().strip().split())\nma...	[]	
6	bc937cacda9ebff9ec0b7f00f0f97508	Sereja has an $n \times m$ rectangular table a, each ...	2200.0	#!/usr/bin/python\n\nimport sys\nfrom math imp...	[]	
...
4970	e745f777e5676c2629f806e09f83a739	Vitaly enrolled in the course Advanced Useless...	2200.0	import os\n\nimport sys\n\nfrom io import Byte...	[]	
4971	a04cfd22f90b2b87f7c5b48dfe3de873	Consider a hallway, which can be represented a...	2400.0	import sys\n\ninput = lambda: sys.stdin.readli...	[]	

```
data[data.tags.apply(lambda x: any(each!=each for each in x))]
```

```
src_uid problem_description difficulty source_code tags
```

```
data.loc[data.tags.apply(lambda x: len(x) == 0), 'tags'] = None
```



```
data = data.dropna()
```

```
data.shape, data.src_uid.nunique()
```

```
((2656, 5), 2656)
```

```
data
```

	src_uid	problem_description	difficulty	source_code	tags
0	9a5bd9f937da55c3d26d5ecde6e50280	Ayrat has number n, represented as it's prime ...	2000.0	from collections import Counter\nfrom operator...	[number theory, math]
1	8ce1ba0a98031c1bc28f53c11905391c	Petya loves lucky numbers. Everybody knows tha...	1500.0	# 447 or 477 start with odd\n\n\n, k = map(in...	[strings]
7	16d4035b138137bbad247ccd5e560051	One day Polycarp published a funny picture in ...	1200.0	from collections import defaultdict\n\nn=int(i...	[graphs, trees]
8	d3b9ffa76436b957ca959cf9204f9873	You've got another problem dealing with arrays...	1700.0	n, k = map(int, input().split())\na = sorted(m...	[math]
10	6f0d3a7971ffc2571838ecd8bf14238d	You are given a grid with \$\$\$n\$\$\$ rows and \$\$\$...	800.0	test = int(input())\r\nsteps = 0\r\nans = []\r...	[math]
...
4976	f26a979dc042ec9564cfecce29e5a1cf	Egor is a famous Russian singer, rapper, actor...	2500.0	import sys\ninput = sys.stdin.readline\nfrom c...	[graphs]
4978	816a82bee65cf79ba8e4d61babcd0301	You are given a tuple generator \$\$\$f(k) = (...	2900.0	import sys\nimport time\nmod=1000000007\ndef s...	[number theory]

Next steps:

[Generate code with data](#)
[New interactive sheet](#)

Modelling Approaches

```
X_data = data.drop(columns=['tags'])
y_data = data.tags
X_train, X_test, y_train, y_test = train_test_split(X_data, y_data, test_size=0.25, random_state=42)
### Here what I usually do in in imbalanced classification problems
### is a "stratified" train_test_split, here we have a MultiClass-Multilabel
### classification problem, so the stratification proposed by sklearn can't handle the labels
### What I do in this case is to use 'iterative_train_test_split' from 'skmultilearn',
### some data transformation should be done first, for now I will not use iterative_train_test_split
### because for the first modelling approach I'm going to use an LLM with one/few shot learning.
```

```
X_data
```

	src_uid	problem_description	difficulty	source_code	
0	9a5bd9f937da55c3d26d5ecde6e50280	Ayrat has number n, represented as it's prime ...	2000.0	from collections import Counter\nfrom operator...	
1	8ce1ba0a98031c1bc28f53c11905391c	Petya loves lucky numbers. Everybody knows tha...	1500.0	# 447 or 477 start with odd\n\n\n, k = map(in...	
7	16d4035b138137bbad247ccd5e560051	One day Polycarp published a funny picture in ...	1200.0	from collections import defaultdict\n\nn=int(i...	
8	d3b9ffa76436b957ca959cf9204f9873	You've got another problem dealing with arrays...	1700.0	n, k = map(int, input().split())\na = sorted(m...	
10	6f0d3a7971ffc2571838ecd8bf14238d	You are given a grid with \$\$\$n\$\$\$ rows and \$\$\$...	800.0	test = int(input())\r\nsteps = 0\r\nans = []\r...	
...
4976	f26a979dc042ec9564cfecce29e5a1cf	Egor is a famous Russian singer, rapper, actor...	2500.0	import sys\ninput = sys.stdin.readline\nfrom c...	
4978	816a82bee65cf79ba8e4d61babcd0301	You are given a tuple generator \$\$\$f(k) = (...	2900.0	import sys\nimport time\nmod=1000000007\ndef s...	

Next steps:

[Generate code with X_data](#)
[New interactive sheet](#)

1st Approach : LLM in one shot.

```
#lets test if our api is functional
client = AzureOpenAI()
```

```

response = client.chat.completions.create(
    model="gpt-4o-mini",
    messages=[
        {"role": "system", "content": "You are a JSON-only response generator."},
        {"role": "user", "content": "yoo"}
    ],
    response_format={"type": "json_object"},
)
print(response.choices[0].message.content)

{
  "response": "Hello! How can I assist you today?"
}

```

The modelling (or the training) part here, is simply to iterate through prompts until finding the best prompt, for this we will use some examples of the training set, once we have the best prompt we will scale and test it on the test set.

```

## 1st iteration

SYSTEM_PROMPT = """
You are an AI assistant that labels CODEFORCES.
"""

USER_PROMPT_1 = """
Based on a PROBLEM_DESCRIPTION of a Codeforces exercises and its respective SOURCE_CODE solution,
You will label the problem by suggesting relevant tags.
The suggested tags have to be drawn from this list : ['math', 'graphs', 'strings', 'number theory', 'trees', 'geometry', 'games']
Don't create any other tags.

You need to understand the problem, the solution, the meaning of each tag and only output the best tag or tags that match the problem.
Output a json list of tags separated by commas.

PROBLEM_DESCRIPTION : {problem_description}
SOURCE_CODE : {source_code}

"""

## 2nd iteration

USER_PROMPT_2 = """
Based on a PROBLEM_DESCRIPTION of a Codeforces exercises and its respective SOURCE_CODE solution,
You will label the problem by suggesting relevant tags.
The suggested tags have to be drawn from this list : ['math', 'graphs', 'strings', 'number theory', 'trees', 'geometry', 'games']
No other tags are allowed.

You need to understand the problem, the solution, the meaning of each tag and only output the best tag or tags that match the problem.
Output a json list of tags separated by commas.

Here are short, concise, precise descriptions for each Codeforces problem tag :
math – Problems that require arithmetic, algebra, combinatorics, or general mathematical reasoning and formulas (proofs, computations)
graphs – Problems about vertices and edges: traversals, shortest/longest paths, connectivity, components, cycles, flows, matchings
strings – Problems that focus on sequences of characters: pattern matching, hashing, substring/suffix structures, automata, palindromes
number theory – Problems about integers and their properties: primes, gcd/lcm, modular arithmetic, congruences, divisibility, totient
trees – Problems on acyclic connected graphs (rooted/unrooted): tree traversals, LCA, centroid/decomposition, rerooting, subtree
geometry – Problems involving points, lines, polygons, distances, angles, vectors, intersections, areas, and computational geometry
games – Problems about combinatorial games and strategies: impartial/partisan games, winning positions, Sprague-Grundy theorem, minimax
probabilities – Problems using chance and expectation: probability calculations, expected value, linearity of expectation, distributions

Make sure the JSON is VALID.

Here are a few examples :
PROBLEM_DESCRIPTION : {problem_description_1}
SOURCE_CODE : {source_code_1}
TAGS : {tags_1}

PROBLEM_DESCRIPTION : {problem_description_2}
SOURCE_CODE : {source_code_2}
TAGS : {tags_2}

Label this problem :
PROBLEM_DESCRIPTION : {problem_description}
SOURCE_CODE : {source_code}

"""

```

First iteration : llm prompt that describes the task step by step and include rules to be respected by the llm.

Second iteration : prompt that describes the task step by step, include rules to be respected by the llm but also provides some examples of problems from the training set with their respective labels.

```
def tag_with_llm(client=client, user_prompt=""):
    sleep(1)
    response = client.chat.completions.create(
        model="gpt-4o-mini",
        messages=[
            {"role": "system", "content": SYSTEM_PROMPT},
            {"role": "user", "content": user_prompt}
        ],
        response_format={"type": "json_object"},
    )
    return json.loads(response.choices[0].message.content).get('tags')
```

```
# The training phase here is to test the prompt with a few examples
# for speed purposes and improve the prompt each time until finding
# the best prompt (lets try with 1/15 of the train set, less than 200 examples)
X_prompt = X_train[:len(X_train)//15]
y_prompt = y_train[:len(y_train)//15]
```

```
X_prompt.shape , y_prompt.shape
```

```
((132, 4), (132,))
```

▼ First prompt

```
#one example testing
tag_with_llm(user_prompt=USER_PROMPT_1.format(problem_description=X_prompt.iloc[0].problem_description, source_code=X_prompt.il
['math']
```

```
y_prompt.iloc[0]
```

```
['math']
```

```
y_prompt_predicted = []
for i in range(X_prompt.shape[0]):
    y_prompt_predicted.append(tag_with_llm(user_prompt=USER_PROMPT_1.format(problem_description=X_prompt.iloc[i].problem_descript
```

```
sum([e==None for e in y_prompt_predicted])
```

```
6
```

```
# Re run the model on examples that where not processed correctly by the llm
for i in range(len(y_prompt_predicted)):
    if y_prompt_predicted[i] == None:
        y_prompt_predicted[i] = tag_with_llm(user_prompt=USER_PROMPT_1.format(problem_description=X_prompt.iloc[i].problem_descript
```

```
sum([e==None for e in y_prompt_predicted])
```

```
0
```

```
### We will use 'weighted f1 score' for evaluation, in order not to penalize the model for the smaller classes
mlb = MultiLabelBinarizer()
y_prompt_encoded = mlb.fit_transform(y_prompt)
y_prompt_predicted_encoded = mlb.fit_transform(y_prompt_predicted)

weighted = f1_score(y_prompt_encoded, y_prompt_predicted_encoded, average='weighted')
```

```
print('First prompt f1 score :',weighted)
```

```
First prompt f1 score : 0.6890580209002773
```

▼ Second prompt

```
#one example testing
tag_with_llm(user_prompt=USER_PROMPT_2.format(
    problem_description_1=X_prompt.iloc[-1].problem_description,
    source_code_1=X_prompt.iloc[-1].source_code,
```

```
tags_1=y_prompt.iloc[-1],

problem_description_2=X_prompt.iloc[-2].problem_description,
source_code_2=X_prompt.iloc[-2].source_code,
tags_2=y_prompt.iloc[-2],

problem_description=X_prompt.iloc[0].problem_description,
source_code=X_prompt.iloc[0].source_code))
```

```
['math']
```

```
y_prompt_predicted = []
for i in range(X_prompt.shape[0]):
    y_prompt_predicted.append(tag_with_llm(user_prompt=USER_PROMPT_2.format(
        problem_description_1=X_prompt.iloc[-1].problem_description,
        source_code_1=X_prompt.iloc[-1].source_code,
        tags_1=y_prompt.iloc[-1],

        problem_description_2=X_prompt.iloc[-2].problem_description,
        source_code_2=X_prompt.iloc[-2].source_code,
        tags_2=y_prompt.iloc[-2],

        problem_description=X_prompt.iloc[i].problem_description, source_code=X_prompt.iloc[i].source_code)))

while sum([e==None for e in y_prompt_predicted])>0:
    for i in range(len(y_prompt_predicted)):
        if not y_prompt_predicted[i]:
            y_prompt_predicted[i] = tag_with_llm(user_prompt=USER_PROMPT_2.format(
                problem_description_1=X_prompt.iloc[-1].problem_description,
                source_code_1=X_prompt.iloc[-1].source_code,
                tags_1=y_prompt.iloc[-1],

                problem_description_2=X_prompt.iloc[-2].problem_description,
                source_code_2=X_prompt.iloc[-2].source_code,
                tags_2=y_prompt.iloc[-2],

                problem_description=X_prompt.iloc[i].problem_description, source_code=X_prompt.iloc[i].source_code)))
```

```
mlb = MultiLabelBinarizer()
y_prompt_encoded = mlb.fit_transform(y_prompt)
y_prompt_predicted_encoded = mlb.fit_transform(y_prompt_predicted)

weighted = f1_score(y_prompt_encoded, y_prompt_predicted_encoded, average='weighted')
```

```
print('Second prompt f1 score :',weighted)
```

```
Second prompt f1 score : 0.6955657774072096
```

```
print(classification_report(y_prompt_encoded, y_prompt_predicted_encoded))
```

	precision	recall	f1-score	support
0	0.30	0.75	0.43	4
1	0.71	0.83	0.77	12
2	0.84	0.59	0.70	27
3	0.78	0.58	0.67	65
4	0.60	0.60	0.60	15
5	0.75	0.60	0.67	5
6	0.81	0.81	0.81	26
7	0.91	0.62	0.74	16
micro avg	0.74	0.65	0.69	170
macro avg	0.71	0.67	0.67	170
weighted avg	0.77	0.65	0.70	170
samples avg	0.76	0.69	0.70	170

We see a small improvement in the f1 score, We choose the second prompt as our model, and we will run it through the test set to evaluate its performance

```
### RUN ON TEST SET
```

```
y_prompt_predicted = []
for i in range(X_test.shape[0]):
    y_prompt_predicted.append(tag_with_llm(user_prompt=USER_PROMPT_2.format(
        problem_description_1=X_prompt.iloc[-1].problem_description,
        source_code_1=X_prompt.iloc[-1].source_code,
```

```

tags_1=y_prompt.iloc[-1],

problem_description_2=X_prompt.iloc[-2].problem_description,
source_code_2=X_prompt.iloc[-2].source_code,
tags_2=y_prompt.iloc[-2],

    problem_description=X_test.iloc[i].problem_description, source_code=X_test.iloc[i].source_code)))

while sum([e==None for e in y_prompt_predicted])>0:
    for i in range(len(y_prompt_predicted)):
        if y_prompt_predicted[i] == None:
            y_prompt_predicted[i] = tag_with_llm(user_prompt=USER_PROMPT_2.format(
                problem_description_1=X_prompt.iloc[-1].problem_description,
                source_code_1=X_prompt.iloc[-1].source_code,
                tags_1=y_prompt.iloc[-1],

                problem_description_2=X_prompt.iloc[-2].problem_description,
                source_code_2=X_prompt.iloc[-2].source_code,
                tags_2=y_prompt.iloc[-2],

                problem_description=X_test.iloc[i].problem_description, source_code=X_test.iloc[i].source_code)
            )

```

```

mlb = MultiLabelBinarizer()
y_prompt_encoded = mlb.fit_transform(y_test)
y_prompt_predicted_encoded = mlb.fit_transform(y_prompt_predicted)

weighted = f1_score(y_prompt_encoded, y_prompt_predicted_encoded, average='weighted')

```

```
print('F1 Score on test set: ',weighted)
```

```
F1 Score on test set: 0.711880222655802
```

```
print(classification_report(y_prompt_encoded, y_prompt_predicted_encoded))
```

	precision	recall	f1-score	support
0	0.31	1.00	0.47	19
1	0.63	0.88	0.73	33
2	0.88	0.62	0.73	156
3	0.82	0.56	0.67	323
4	0.62	0.74	0.67	86
5	0.80	0.70	0.74	23
6	0.75	0.86	0.80	107
7	0.88	0.74	0.80	85
micro avg	0.74	0.68	0.71	832
macro avg	0.71	0.76	0.70	832
weighted avg	0.79	0.68	0.71	832
samples avg	0.77	0.72	0.72	832

2nd Approach : BERT Classifier

```

labels_to_ids = {j:i for i,j in enumerate(focus_tags)}
ids_to_labels = {i:j for i,j in enumerate(focus_tags)}

```

```
labels_to_ids
```

```

{'math': 0,
 'graphs': 1,
 'strings': 2,
 'number theory': 3,
 'trees': 4,
 'geometry': 5,
 'games': 6,
 'probabilities': 7}

```

```

MAX_LEN = 128
TRAIN_BATCH_SIZE = 16
VALID_BATCH_SIZE = 2
EPOCHS = 1
LEARNING_RATE = 1e-05
tokenizer = BertTokenizerFast.from_pretrained('bert-base-uncased')

```

```

/usr/local/lib/python3.12/dist-packages/huggingface_hub/utils/_auth.py:94: UserWarning:
The secret `HF_TOKEN` does not exist in your Colab secrets.
To authenticate with the Hugging Face Hub, create a token in your settings tab (https://huggingface.co/settings/tokens), set it
You will be able to reuse this secret in all of your notebooks.
Please note that authentication is recommended but still optional to access public models or datasets.
warnings.warn(
tokenizer_config.json: 100% 48.0/48.0 [00:00<00:00, 1.08kB/s]
vocab.txt: 100% 232k/232k [00:00<00:00, 3.68MB/s]
tokenizer.json: 100% 466k/466k [00:00<00:00, 17.0MB/s]
config.json: 100% 570/570 [00:00<00:00, 21.4kB/s]

```

```

# Data retriever class, we need to tokenize text and binerize labels
# before feeding the Bert classifier, this class will help us handle
# this part on the fly

```

```

class dataset(Dataset):
    def __init__(self, dataframe, tokenizer, max_len):
        self.len = len(dataframe)
        self.data = dataframe
        self.tokenizer = tokenizer
        self.max_len = max_len

    def __getitem__(self, index):

        sentence = self.data.problem_description[index].strip()
        labels = self.data.tags[index]

        encoding = self.tokenizer(sentence,
                                   return_offsets_mapping=True,
                                   padding='max_length',
                                   truncation=True,
                                   max_length=self.max_len)

        labels_holder = np.array([0 for _ in range(len(labels_to_ids))], dtype=np.float32)
        for label in labels:
            labels_holder[labels_to_ids[label]] = 1
        labels = labels_holder[:]

        item = {key: torch.as_tensor(val) for key, val in encoding.items()}
        item['labels'] = torch.as_tensor(labels)
        return item

    def __len__(self):
        return self.len

```

```

# We create an architecture based on a bert model for the embedding phase
# We add a dense layer for classification (we could have added a dropout layer in between)

```

```

class ModelClassifier(pl.LightningModule):

    def __init__(self, n_classes: int, model_name: str):
        super().__init__()
        self.bert = BertModel.from_pretrained(model_name, return_dict=True)
        self.classifier = nn.Linear(self.bert.config.hidden_size, n_classes)
        self.criterion = nn.BCELoss()

    def forward(self, input_ids, attention_mask, labels=None):
        output = self.bert(input_ids, attention_mask=attention_mask)
        output = self.classifier(output.pooler_output)
        output = torch.sigmoid(output)
        loss = 0
        if labels is not None:
            loss = self.criterion(output, labels)
        return loss, output

    def training_step(self, batch, batch_idx):
        input_ids = batch["input_ids"]
        attention_mask = batch["attention_mask"]
        labels = batch["labels"]
        loss, outputs = self(input_ids, attention_mask, labels)
        self.log("train_loss", loss, prog_bar=True, logger=True)
        return {"loss": loss, "predictions": outputs, "labels": labels}

    def validation_step(self, batch, batch_idx):
        input_ids = batch["input_ids"]
        attention_mask = batch["attention_mask"]
        labels = batch["labels"]
        loss, outputs = self(input_ids, attention_mask, labels)

```

```

self.log("val_loss", loss, prog_bar=True, logger=True)
return loss

def configure_optimizers(self):
    optimizer = torch.optim.AdamW(params=self.parameters(), lr=LEARNING_RATE)
    scheduler = get_linear_schedule_with_warmup(
        optimizer,
        num_warmup_steps=20,
        num_training_steps=100
    )
    return dict(
        optimizer=optimizer,
        lr_scheduler=dict(
            scheduler=scheduler,
            interval='step'
        )
    )

```

```

train_dataset = data[data.index.isin(X_train.index)]
test_dataset = data.drop(train_dataset.index).reset_index(drop=True)
train_dataset = train_dataset.reset_index(drop=True)

```

```

print("FULL Dataset: {}".format(data.shape))
print("TRAIN Dataset: {}".format(train_dataset.shape))
print("TEST Dataset: {}".format(test_dataset.shape))

```

```

training_set = dataset(train_dataset, tokenizer, MAX_LEN)
testing_set = dataset(test_dataset, tokenizer, MAX_LEN)

```

```

FULL Dataset: (2656, 5)
TRAIN Dataset: (1992, 5)
TEST Dataset: (664, 5)

```

```

train_params = {'batch_size': TRAIN_BATCH_SIZE,
                'shuffle': True,
                'num_workers': 0
                }

```

```

test_params = {'batch_size': VALID_BATCH_SIZE,
               'shuffle': True,
               'num_workers': 0
               }

```

```

training_loader = DataLoader(training_set, **train_params)
testing_loader = DataLoader(testing_set, **test_params)

```

```

model = ModelClassifier(len(labels_to_ids), 'bert-base-uncased')
#trainer = pl.Trainer(accelerator="cpu", devices=1, max_epochs=3)
#trainer.fit(model, training_loader)

```

```
model.safetensors: 100%
```

```
440M/440M [00:21<00:00, 22.2MB/s]
```

```

model = ModelClassifier(len(labels_to_ids), 'bert-base-uncased')
state_dict = torch.load("model_v1", weights_only=True, map_location=device)
model.load_state_dict(state_dict)

```

```
<All keys matched successfully>
```

```

THRESHOLD = 0.5
y_pred = []
y_true = []

model.eval()
with torch.no_grad():
    for idx, batch in enumerate(testing_loader):

        ids = batch['input_ids']
        mask = batch['attention_mask']
        labels = batch['labels']

        loss, predictions = model(input_ids=ids, attention_mask=mask, labels=labels)
        predictions = (predictions > THRESHOLD).long()

        y_pred.append(predictions)
        y_true.append(labels)

```

```
metric = MulticlassF1Score(num_classes=len(labels_to_ids), average="weighted")
f1_tensor = metric(predictions, labels)

f1_value = f1_tensor.item() if isinstance(f1_tensor, torch.Tensor) else float(f1_tensor)
print("Weighted multiclass F1:", f1_value)
```

Weighted multiclass F1: 0.9281609058380127

The bert based model clearly outperforms the one shot LLM approach. This is not surprising since the BERT is trained but the LLM was not.

We will use the bert model for the cli module.

```
!python problem_tagger.py --test true
```

```
2025-10-28 15:34:06.603398: E external/local_xla/xla/stream_executor/cuda/cuda_fft.cc:467] Unable to register cuFFT factory: Att
WARNING: All log messages before absl::InitializeLog() is called are written to STDERR
E0000 00:00:1761665646.628375 41830 cuda_dnn.cc:8579] Unable to register cudNN factory: Attempting to register factory for plu
E0000 00:00:1761665646.635450 41830 cuda_blas.cc:1407] Unable to register cuBLAS factory: Attempting to register factory for p
W0000 00:00:1761665646.653355 41830 computation_placer.cc:177] computation placer already registered. Please check linkage and
W0000 00:00:1761665646.653402 41830 computation_placer.cc:177] computation placer already registered. Please check linkage and
W0000 00:00:1761665646.653407 41830 computation_placer.cc:177] computation placer already registered. Please check linkage and
W0000 00:00:1761665646.653411 41830 computation_placer.cc:177] computation placer already registered. Please check linkage and
{'TRUE_LABELS': ['math'], 'PREDICTED_LABELS': ['math']}
```

Improvements

This work was done in 4h approximatively. If more time is to be invested in the project we can explore many approaches/techniques that could improve performance.

DATASET improvements :

- We have an imbalanced dataset, that's the first thing to be approached. Historically techniques that uses synthetic data augmentation (like SMOTE) are the first thing to be tested. But here we are lucky enough to have a real world dataset publically available. So what we can do is simply augment the very small classes with real world codeforces problems. ATTENTION : This does not mean creating a dataset with all classes equally presented, because that's not what the real world is (real world : codeforces problems dataset). So we will trv to have a balance between minimum necessarv number of samples for the model to learn and real